


1



2

Outline

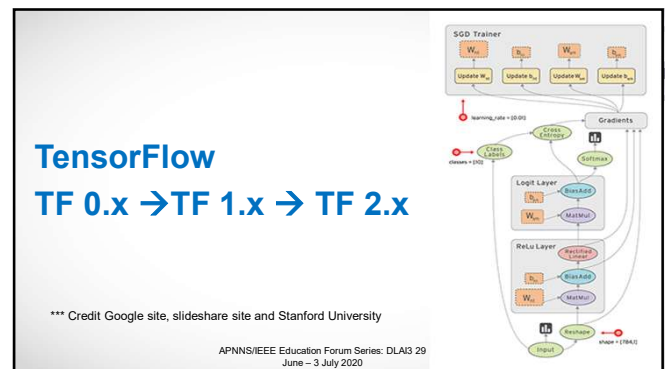
- Creating DL models
 - TensorFlow, and Keras
 - Sequential
 - Loss function and Training
 - Practical: Linear Classifier, MLP, and CNN

Acknowledgement: images and materials are from

- tensorflow.org
- <https://codealabs.developers.google.com/codealabs/cloud-tensorflow-mnist/#0> (by Martin Gorner)

APNNS/IEEE Education Forum Series: DLA13 29 June – 3 July 2020

3



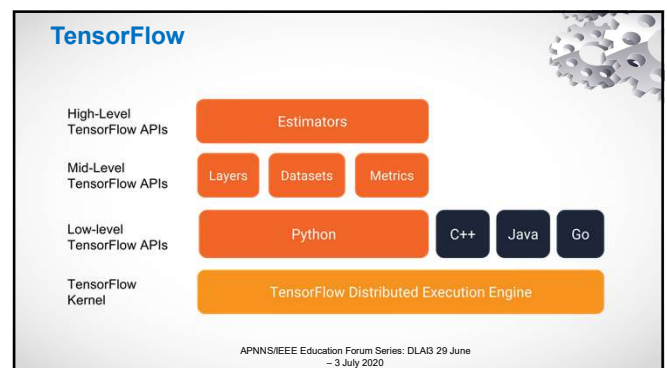
4

Tensorflow

- Starting in 2011, Google Brain built **DistBelief** as a proprietary machine learning system based on deep learning neural networks.
- **TensorFlow** is Google Brain's second-generation system. Version 1.0.0 was released on February 11, 2017 (**version 2.0 was released in October 2019**).
- TensorFlow computations are expressed as stateful **dataflow graphs**. The name TensorFlow derives from the operations that such neural networks perform on **multidimensional data arrays**, which are referred to as **tensors**.
- TensorFlow provides stable Python (for version 3.7 across all platforms) and C APIs; and without API backwards compatibility guarantee: C++, Go, Java, JavaScript and Swift (early release).

APNNS/IEEE Education Forum Series: DLA13 29 June – 3 July 2020

5




6

Graphs and Computations

TensorFlow Graph Execution separates definition of computations from their execution

1. Phase 1: assemble a graph
2. Phase 2: use a session (TF1.x) to execute operations in the graph.



APNNS/IEEE Education Forum Series: DLA3 29
June - 3 July 2020

7

Benefits of Graphs

1. Save computation. Only run subgraphs that lead to the values you want to fetch.
2. Break computation into small, differential pieces to facilitate auto-differentiation
3. Facilitate distributed computation, spread the work across multiple CPUs, GPUs, TPUs, or other devices
4. Many common machine learning models are taught and visualized as directed graphs

APNNS/IEEE Education Forum Series: DLA3 29
June - 3 July 2020

8

Tensors

- An n-dimensional array
 - 0-d tensor: scalar (number)
 - 1-d tensor: vector
 - 2-d tensor: matrix
 - and so on

APNNS/IEEE Education Forum Series: DLA3 29
June - 3 July 2020

9

Data Flow Graph

```
import tensorflow as tf
a = tf.add(3, 5)
```



TF automatically names the nodes when you don't explicitly name them.

x = 3
y = 5

APNNS/IEEE Education Forum Series: DLA3 29
June - 3 July 2020

10

Data Flow Graph

```
import tensorflow as tf
a = tf.add(3, 5)
print(a)

>> tf.Tensor(8, shape=(), dtype=int32)
```



APNNS/IEEE Education Forum Series: DLA3 29 June -
3 July 2020

11

Tensor Constructors

```
tf.zeros([2, 3], tf.int32) ==> [[0, 0, 0], [0, 0, 0]]

# input_tensor is [[0, 1], [2, 3], [4, 5]]
tf.zeros_like(input_tensor) ==> [[0, 0], [0, 0], [0, 0]]

tf.fill([2, 3], 8) ==> [[8, 8, 8], [8, 8, 8]]

tf.linspace(start, stop, num, name=None)
tf.linspace(10.0, 13.0, 4) ==> [10. 11. 12. 13.]

tf.range(start, limit=None, delta=1, dtype=None, name='range')
tf.range(3, 18, 3) ==> [3 6 9 12 15]
tf.range(5) ==> [0 1 2 3 4]
```


APNNS/IEEE Education Forum Series: DLA3 29
June - 3 July 2020

12

Practical: TensorFlow

*** Credit Google site, slideshare site and Stanford University

APNNS/IEEE Education Forum Series: DLA3 29 June - 3 July 2020



13

Example 1: Quadratic Equation

Creating a Simple TensorFlow Program

Let's write a tensorflow graph to compute roots of a given quadratic equation

$$y = ax^2 + bx + c$$

$$e.g., y = 2x^2 - 3x - 1$$

```
a = tf.constant(2)
b = tf.constant(-3)
c = tf.constant(-1)
```

APNNS/IEEE Education Forum Series: DLA3 29 June - 3 July 2020

14


Example 2: Perceptron

Output

$$\hat{y} = g \left(\sum_{i=1}^m x_i w_i \right)$$

Linear combination of inputs

Non-linear activation function




APNNS/IEEE Education Forum Series: DLA3 29 June - 3 July 2020 © MIT 6.S191: Introduction to Deep Learning introtodeeplearning.com

15

We have: $w_0 = 1$ and $\mathbf{W} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$

$$\begin{aligned} \hat{y} &= g(w_0 + \mathbf{X}^T \mathbf{W}) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ &= g(1 + 3x_1 - 2x_2) \end{aligned}$$





APNNS/IEEE Education Forum Series: DLA3 29 June - 3 July 2020 © MIT 6.S191: Introduction to Deep Learning introtodeeplearning.com

16

We have: $w_0 = 1$ and $\mathbf{W} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$

$$\begin{aligned} \hat{y} &= g(w_0 + \mathbf{X}^T \mathbf{W}) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ &= g(1 + 3x_1 - 2x_2) \end{aligned}$$



APNNS/IEEE Education Forum Series: DLA3 29 June - 3 July 2020 © MIT 6.S191: Introduction to Deep Learning introtodeeplearning.com

17

Keras

- Keras is an open-source neural-network library written in Python.
- It is capable of running on top of TensorFlow, Microsoft Cognitive Toolkit, R, Theano, or PlaidML.
- It is designed to enable fast experimentation with deep neural networks, it focuses on being user-friendly, modular, and extensible. Its primary author and maintainer is François Chollet, a Google engineer.



APNNS/IEEE Education Forum Series: DLA3 29 June - 3 July 2020

18

keras & tensorflow.keras

- TensorFlow's implementation of the Keras API is accessed via `tf.keras`.
- This is a high-level API to build and train DL models.
- `'tf.keras'` makes TensorFlow easier to use without sacrificing flexibility and performance.

APNNS/IEEE Education Forum Series: DLA3 29
June – 3 July 2020

19

tensorflow.keras.Sequential

```
tf.keras.Sequential(layers=None, name=None)

model = tf.keras.Sequential()
model.add(tf.keras.layers.Dense(8, input_shape=(16,)))
...
model.add(tf.keras.layers.Dense(4))
```

APNNS/IEEE Education Forum Series: DLA3 29
June – 3 July 2020

20

tensorflow.keras.Sequential

```
compile(
    optimizer='rmsprop', loss=None, metrics=None, loss_weights=None,
    sample_weight_mode=None, weighted_metrics=None, **kwargs
)

fit(
    x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
    validation_split=0.0, validation_data=None, shuffle=True, class_weight=None,
    sample_weight=None, initial_epoch=0, steps_per_epoch=None,
    validation_steps=None, validation_batch_size=None, validation_freq=1,
    max_queue_size=10, workers=1, use_multiprocessing=False
)
```

APNNS/IEEE Education Forum Series: DLA3 29
June – 3 July 2020

21

tensorflow.keras.Sequential

```
evaluate(
    x=None, y=None, batch_size=None, verbose=1, sample_weight=None,
    steps=None, callbacks=None, max_queue_size=10, workers=1,
    use_multiprocessing=False, return_dict=False
)

predict(
    x, batch_size=None, verbose=0, steps=None, callbacks=None,
    max_queue_size=10, workers=1, use_multiprocessing=False
)
```

APNNS/IEEE Education Forum Series: DLA3 29
June – 3 July 2020

22

Building a Model in Keras using Sequential

```
model = tf.keras.Sequential()

# Adds a densely-connected layer with 64 units to the model:
model.add(tf.keras.layers.Flatten(input_shape=[28,28,1]))

# Add another:
model.add(tf.keras.layers.Dense(64, activation='relu'))

# Add a softmax layer with 10 output units:
model.add(tf.keras.layers.Dense(10, activation='softmax'))
```

APNNS/IEEE Education Forum Series: DLA3 29
June – 3 July 2020

23

Example of Layers in tf.keras.layers

```
# Create a sigmoid layer:
layers.Dense(64, activation='tf.sigmoid')

# A linear layer with L2 regularization of factor 0.01 applied to the bias vector:
layers.Dense(64, bias_regularizer=tf.keras.regularizers.l2(0.01))

# A 2D convolutional layer
layers.Conv2D(kernel_size=6, filters=24, strides=2, activation='relu')

# A polling layer
layers.MaxPooling2D(pool_size=(2,2), stride=2, padding='same')
```

APNNS/IEEE Education Forum Series: DLA3 29
June – 3 July 2020

24

Build and Compile a Model

```
model = tf.keras.Sequential([
    layers.Dense(64, activation='relu', input_shape=(32,)),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')])

model.compile(optimizer=tf.train.AdamOptimizer(0.001),
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

APNNS/IEEE Education Forum Series: DLA3 29
June - 3 July 2020

25

Train, Save and Load Model

```
# Train model
model.fit(data, labels, batch_size=32, epochs=5)


# Save entire model to a HDF5 file
model.save('my_model.h5')

# Recreate the exact same model, including weights and optimizer.
model = tf.keras.models.load_model('my_model.h5')
```

APNNS/IEEE Education Forum Series: DLA3 29
June - 3 July 2020

26


Loss Function



APNNS/IEEE Education Forum Series: DLA3 29 June
© MIT 6.S191: Introduction to Deep Learning
- 3 July 2020

27

Mean-Square Error Loss




```
loss = tf.reduce_mean( tf.square( tf.subtract( model.y, model.pred ) ) )
```

APNNS/IEEE Education Forum Series: DLA3 29 June
© MIT 6.S191: Introduction to Deep Learning
- 3 July 2020

28

Binary Cross Entropy Loss



```
loss = tf.reduce_mean( tf.nn.softmax_cross_entropy_with_logits(model.y, model.pred) )
```

APNNS/IEEE Education Forum Series: DLA3 29 June
© MIT 6.S191: Introduction to Deep Learning
- 3 July 2020

29

Gradient Descent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights

```
weights = tf.random_normal(shape, stddev=stddev)
```

```
grads = tf.gradients(losses, weights)
```


```
weights_new = weights.assign(weights - lr * grads)
```

APNNS/IEEE Education Forum Series: DLA3 29 June
© MIT 6.S191: Introduction to Deep Learning
- 3 July 2020

30

Gradient Descent


Randomly pick an initial (w_0, w_1)



© MIT 6.S191: Introduction to Deep Learning
APNNS/IEEE Education Forum Series: DLA13 29 Jun 2020
introtodeeplearning.com

31


Gradient Descent



© MIT 6.S191: Introduction to Deep Learning
APNNS/IEEE Education Forum Series: DLA13 29 Jun 2020
introtodeeplearning.com

32

Gradient Descent




© MIT 6.S191: Introduction to Deep Learning
APNNS/IEEE Education Forum Series: DLA13 29 Jun 2020
introtodeeplearning.com

33

Gradient Descent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(W)}{\partial W}$
4. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights




© MIT 6.S191: Introduction to Deep Learning
APNNS/IEEE Education Forum Series: DLA13 29 Jun 2020
introtodeeplearning.com

34

Stochastic Gradient Descent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick single data point i
4. Compute gradient, $\frac{\partial J_i(W)}{\partial W}$
5. Update weights, $W \leftarrow W - \eta \frac{\partial J_i(W)}{\partial W}$
6. Return weights



Easy to compute but
very noisy
(stochastic)!


© MIT 6.S191: Introduction to Deep Learning
APNNS/IEEE Education Forum Series: DLA13 29 Jun 2020
introtodeeplearning.com

35

Minibatch Gradient Descent

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick batch of B data points
4. Compute gradient, $\frac{\partial J(W)}{\partial W} = \frac{1}{B} \sum_{k=1}^B \frac{\partial J_k(W)}{\partial W}$
5. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
6. Return weights




© MIT 6.S191: Introduction to Deep Learning
APNNS/IEEE Education Forum Series: DLA13 29 Jun 2020
introtodeeplearning.com

36

Practical: TensorFlow


*** Credit Google site, slideshare site and Stanford University

APNNS/IEEE Education Forum Series: DLAI3 29 June - 3 July 2020



37

Example 3: Multi-Layer Perceptron




$$z_2 = w_{0,2}^{(1)} + \sum_{j=1}^m x_j w_{j,2}^{(1)} \\ = w_{0,2}^{(1)} + x_1 w_{1,2}^{(1)} + x_2 w_{2,2}^{(1)} + x_m w_{m,2}^{(1)}$$

© MIT 6.S191: Introduction to Deep Learning
APNNS/IEEE Education Forum Series: DLAI3 29 June - 3 July 2020

38


Example 4: Linear Softmax Classifier



© MIT 6.S191: Introduction to Deep Learning
APNNS/IEEE Education Forum Series: DLAI3 29 June - 3 July 2020

39

Example 4: Linear Softmax Classifier

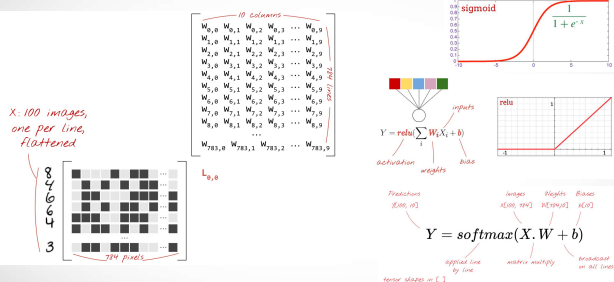


© MIT 6.S191: Introduction to Deep Learning
introtodeeplearning.com

APNNS/IEEE Education Forum Series: DLAI3 29 June - 3 July 2020

40


Example 4: Linear Softmax Classifier



APNNS/IEEE Education Forum Series: DLAI3 29 June - 3 July 2020


41

Example 4: Linear Softmax Classifier



APNNS/IEEE Education Forum Series: DLAI3 29 June - 3 July 2020

42



43

Terminology Covered So Far (1)

- batch or mini-batch: training is always performed on batches of training data and labels. Doing so helps the algorithm converge. The "batch" dimension is typically the first dimension of data tensors. For example a tensor of shape [100, 192, 192, 3] contains 100 images of 192x192 pixels with three values per pixel (RGB).
- cross-entropy loss: a special loss function often used in classifiers.
- dense layer: a layer of neurons where each neuron is connected to all the neurons in the previous layer.

APNNS/IEEE Education Forum Series: DLAI3 29
June - 3 July 2020

44

Terminology Covered So Far (2)

- features: the inputs of a neural network are sometimes called "features". The art of figuring out which parts of a dataset (or combinations of parts) to feed into a neural network to get good predictions is called "feature engineering".
- labels: another name for "classes" or correct answers in a supervised classification problem
- learning rate: fraction of the gradient by which weights and biases are updated at each iteration of the training loop.

APNNS/IEEE Education Forum Series: DLAI3 29
June - 3 July 2020

45

Terminology Covered So Far (3)

- logits: the outputs of a layer of neurons before the activation function is applied are called "logits". The term comes from the "logistic function" a.k.a. the "sigmoid function" which used to be the most popular activation function. "Neuron outputs before logistic function" was shortened to "logits".
- loss: the error function comparing neural network outputs to the correct answers
- neuron: computes the weighted sum of its inputs, adds a bias and feeds the result through an activation function.
- one-hot encoding: class 3 out of 5 is encoded as a vector of 5 elements, all zeros except the 3rd one which is 1.

APNNS/IEEE Education Forum Series: DLAI3 29
June - 3 July 2020

46

Terminology Covered So Far (4)


- relu: rectified linear unit. A popular activation function for neurons.
- sigmoid: another activation function that used to be popular and is still useful in special cases.
- softmax: a special activation function that acts on a vector, increases the difference between the largest component and all others, and also normalizes the vector to have a sum of 1 so that it can be interpreted as a vector of probabilities. Used as the last step in classifiers.
- tensor: A "tensor" is like a matrix but with an arbitrary number of dimensions. A 1-dimensional tensor is a vector. A 2-dimensions tensor is a matrix. And then you can have tensors with 3, 4, 5 or more dimensions.

APNNS/IEEE Education Forum Series: DLAI3 29
June - 3 July 2020

47

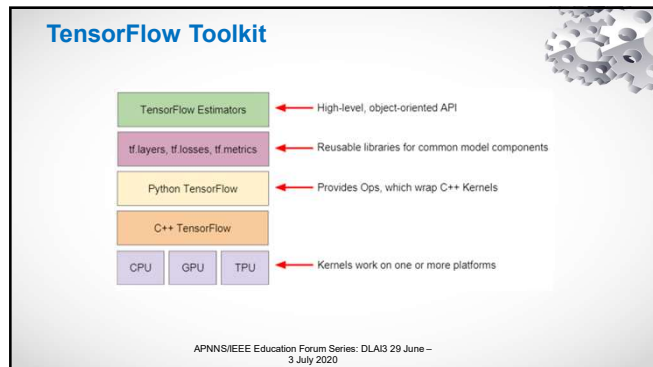
Recap: TensorFlow

*** Credit Google site, slideshare site and Stanford University

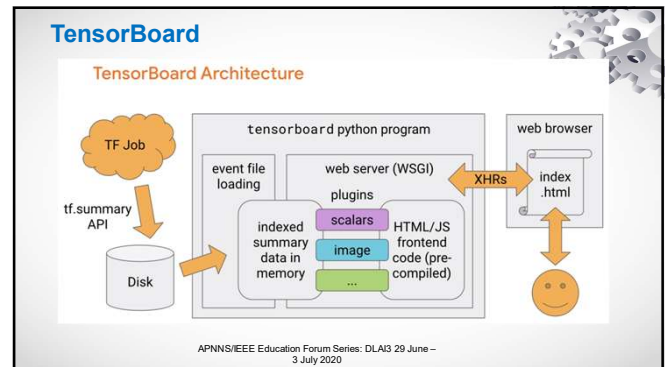


APNNS/IEEE Education Forum Series: DLAI3 29
June - 3 July 2020


48



49



50



51



52